

# Computer Music: Languages and Systems - Homework 2

Francesco Colotti (*francesco.colotti@mail.polimi.it*)  
Gioele Fortugno (*gioele.fortugno@mail.polimi.it*)  
Matteo Gionfriddo (*matteo.gionfriddo@mail.polimi.it*)  
Emanuele Greco (*emanuele2.greco@mail.polimi.it*)

May 30, 2023

In this report, we will briefly analyze *Fire*, a multi-band distortion plugin made with JUCE by jerryuhoo.

## 1 Introduction

*Fire* is practically divided in two tabs, the *Band Effect* part and the *Global Effect* part. Both share a representation of the frequency spectrum of the audio, an Output and Mix knobs and a graph visualizer. The first part lets the user divide the audio spectrum up to a maximum of four bands, for each of which they can set the distortion parameters. The second part instead allows the user to control three filters using a graphic equalizer, which are a low-pass filter, a band-pass filter and a high-pass filter. Finally, in the top region of the plugin an interface is used to do A/B comparisons and preset management.

For our code analysis we will follow a top-down approach: we will first explore the DSP starting from the function `ProcessBlock()` in which the order of signal processing blocks is defined. Then we will concentrate on the function `ProcessOneBand()`, where the single band distortion process is organized. This is followed by a discussion on the `leftChain` and `rightChain` objects, which are crucial elements in the implementation of the *Global Effect* part, then a description of the functions within the `ClippingFunctions.h` file (the core of the distortion processes) and a brief explanation of the function `getStateInformation()`, `setStateInformation()`, important for handling presets and current parameters. We also spend some words on the HQ and the Downsample functionalities of the plugin. Finally there is a general view of The Graphical User Interface.

## 2 Block diagram

### 2.1 Overview

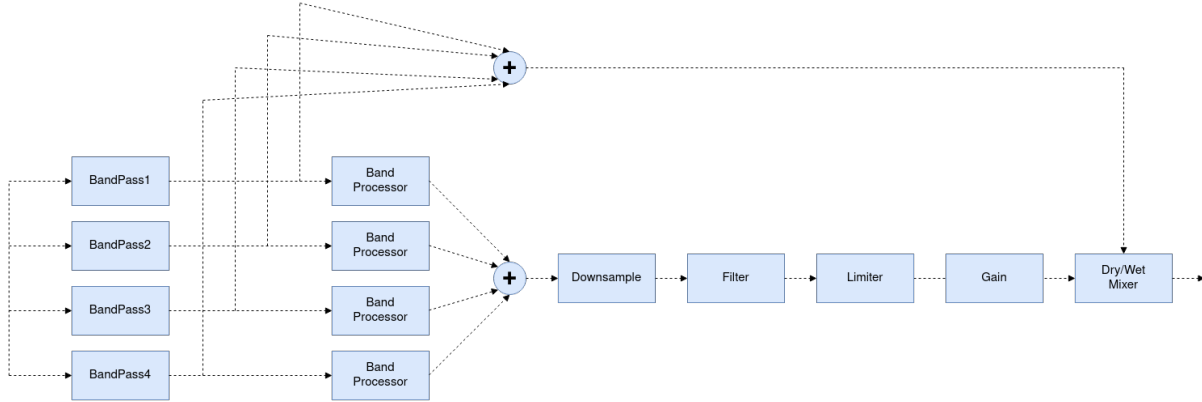


Figure 1: Block Diagram of the behaviour of the plugin

### 2.2 Band Processor




Figure 2: Block Diagram of each Band component

## 3 DSP

### 3.1 ProcessBlock()

As we anticipated, to understand the processing the signal goes through, we need to look at the function `processBlock()` located inside `PluginProcessor.cpp`. In this numerous blocks are defined, though we will analyze only the main ones that provide a rough idea of how the signal is processed. A general block diagram of the signal path is visualized in Figure 2.1.

With the `AudioBlock` class we access the audio data in the buffer to enable audio processing. Following a multiband approach, the audio signal is divided in up to 4 separate frequency bands, with a different processing for each. This frequency band division is applied by properly combining high-pass and low-pass filters. In particular, the first band

from the right and the first band from the left are obtained respectively by means of a high-pass filter and a low-pass filter. Central bands however are obtained by a band-pass filter, made by a sum of an high-pass and a low-pass filter. Obviously two adjacent band division filters share the same cut-off frequency. The `LineNum` variable keeps track of the number of band breakpoints (for example `lineNum = 2` means that the spectrum is divided in 3 bands), while `freqArray` stores the filters cut-off frequencies in an ascending order. By this two elements is ensured that the band division is executed correctly.

The audio buffer is then split into four auxiliary buffers, `mBuffer1`, `mBuffer2`, `mBuffer3`, and `mBuffer4`, one for each band. If the `multibandEnabled` option is true for the current frequency band, the `processOneBand()` function is called to process the buffer. This function does the actual processing of the selected band using the parameters retrieved from the `treeState` object. The values are retrieved using the IDs passed to the function. Finally, the RMS values for each left and right input and output channel are calculated and stored in variables for the current frequency band.

After each band has been processed, the complete output buffer is processed as well. First, the left and right channel of the buffer are processed through the same chain composed of a low-pass, high-pass and peak filter, and then the output is processed through a limiter. Finally the audio block goes through a gain processor. The filters, limiter and gain parameters values are retrieved from the `treeState` parameter tree. Next, the audio signal is mixed with the dry signal according to the `MIX` parameter ID, using a dry-wet mixer. The mixer blends the audio signals while taking into account the plugin's latency. After mixing, the audio buffer is copied to the `mWetBuffer` buffer, which is used to calculate the frequency spectrum of the audio signal through the FFT. The code then computes the RMS values of the left and right channels of the output signal, which are used to drive the VU meter. Finally, the `mDryBuffer` buffer is cleared, preparing it for the next audio block to be processed.

## 3.2 ProcessOneBand()

This function represents the chain through which each band passes and is located inside `PluginProcessor.cpp` file. A general block diagram of the signal path is visualized in Figure 2.2. Other than the actual buffer and the context it takes several arguments, most of which are string ID used to retrieve values from the parameter tree. It calls the following functions, each with the received `bandBuffer` as a parameter.

### 3.2.1 processDistortion()

Here is where the actual distortion is applied. First an `AudioBlock` for the input (from the pointer to the array of channels passed as argument) and for the output are created. Each indication about the plugin settings (e.g. number of band considered, type of distortion ecc...) is retrieved from the `treeState` object and the predefined constants. Based on these values, operations like oversampling or usage of the safe modality are performed, calling the related specified functions defined in the other files. In particular, the pointer `overdrive` to a `BiasProcessor` object is first used to set the distortion amount, the

chosen waveshaping (calling one of the function present in `ClippingFunctions.h`) and the bias and rectification amount. Finally the processing is applied to the current block.

### 3.2.2 WidthProcessor

`WidthProcessor.process()` is applied in case of a stereo signal to make it "wider" or "narrower". It starts by computing the `mid` and `side` component of the buffer, obtained through the sum and the difference of the two signal, divided by the square root of 2. The function then amplify the different components by a factor of `1-width` and `width` respectively before reconstructing the left and right channel from the 2 components.

### 3.2.3 processCompressor() , processGain() , mixDryWet()

. In each case are simply retrieved the associated parameters from the `TreeState` object and, based on them, applied the compression, the gain or the dry-wet process by respectively a `CompressorProcessor`, a `GainProcessor` object or a `DryWetMixer` object. In `mixDryWet()` wet latency is set to a particular value when we are in HQ mode.

## 3.3 leftChain and rightChain

The `leftChain` and `rightChain` are `ProcessorChain` objects that join together the chain of processors that implements the 3-band parametric equalizer, they are instances of `MonoChain` that is defined inside the `FilterControl.h` file. The resonant low-pass(high-pass) filter is implemented by using an IIR high order low-pass(high-pass) filter made with Butterworth method summed with an IIR peak filter, while the band-pass filter needs obviously only an IIR peak filter. As a consequence each process chain is made by 3 peak filter and 2 cut filter. In particular each cut filter is a *ProcessingChain* of 4 filters. Only one of them is active at a given time, depending on the chosen filter slope (12, 24, 36 or 48 dB).

Before the processing happens the function `updateFilter()` is called (inside `processBlock()`). This function gets the filters parameters from *treeState* and calls one helper function for each filter to set the appropriate coefficients. In particular, the corresponding update function of each filter works by first computing the correct coefficients from the parameters obtained through the user input, and then calling the function `UpdateCoefficients()` to set these values inside the filter in the processing chain. In the case of the low-pass filter and the high-pass filter is also updated the related cut filter by the function `UpdateCutFilter()`. After the update the filters are ready to be used and the `process()` function is called on both the chains.

## 3.4 ClippingFunctions.h

In this file all the waveshaping functions designed to distort the signals are defined. The concept behind distortion functions is to alter the incoming audio signal to add harmonics

not present in the original signal. This effect is typically achieved by non-linear saturation of a circuit, which distorts the original audio signal waveform. In this way, richer and more complex sounds can be obtained, ranging from slight saturation to extreme distortion. The distortion functions in this code use various nonlinear saturation algorithms to create distortion effects, and each function has a unique sonic characteristic that can be suitable for different types of sounds and musical styles. The functions take the input signal  $x$ , apply mathematical formulas and return a distorted signal  $x$ . This code is a template `T` that is inserted into the `pluginprocessor.h` and then called from the `pluginprocessor.cpp`. Specifically we have 12 different functions grouped into 3 categories (SoftClipping, HardClipping, Foldback), below we represent some of them:

**CubicSoftClipping:**

$$f(x) = \begin{cases} 1 & \text{if } x > 1 \\ -1 & \text{if } x < -1 \\ \frac{3}{2}x - \frac{x^3}{2} & \text{otherwise} \end{cases}$$

**SausageFattener:**

$$f(z) = \begin{cases} -1 & \text{if } z \leq -1.1 \\ 2.5z^2 + 5.5z + 2.025 & \text{if } -1.1 < z < -0.9 \\ 1 & \text{if } -0.9 \leq z \leq 0.9 \\ -2.5z^2 + 5.5z - 2.025 & \text{if } 0.9 < z < 1.1 \\ 1 & \text{if } z \geq 1.1 \end{cases}$$

with

$$z = 1.1x$$

**LinFoldback:**

$$f(x) = \begin{cases} |(x - 1.4) \bmod(4)| - 2 & \text{if } x > 1 \\ x & \text{otherwise} \end{cases}$$

**LogicClip:**

$$f(x) = \frac{2}{(1 + e^{-2x})} - 1$$

### 3.5 getStateInformation(), setStateInformation()

These methods are used to store and recall parameters in the memory block and are located inside `PluginProcessor.cpp`. This can be done either as raw data, or using the `XML` or `ValueTree` classes as intermediaries to make it easy to save and load complex data.

## 3.6 HQ

The user has the possibility to work at a higher sampling rate to reduce the artifacts introduced by the distortion, at the cost of higher computational complexity. If the HQ mode is enabled, a 4x oversampling is applied inside `processDistortion()` for each block, right before processing it. After the non-linear processing downsampling is performed, also to re-synchronize the output signal with the latency introduced by upsampling.

## 3.7 Downsample

In the Global tab a downsample effect is available. It can be used to obtain a rougher sound and/or to obtain an old school chiptune-ish sound. We chose to include the code for the downsampling of a channel below.

```
//this is repeated for every channel (2 channels for stereo)
//for simplicity we only consider one channel now
//to see the full code check out
auto* channelData = buffer.getWritePointer(channel);
for(int sample=0; sample<buffer.getNumSamples(); ++sample){
    if (rateDivide > 1){
        if (sample % rateDivide != 0)
            channelData[sample] =
                channelData[sample - sample % rateDivide];
    }
}
```

# 4 The Graphical User Interface

For our analysis of the implementation of the GUI we will consider the plugin as made by three panels: Top Panel (the one that allows preset management), Spectrogram Panel (the one in which the audio frequency spectrum is displayed) and Control Panel (the lower part with knobs and various visualizations). Some screenshots of the GUI are available in the *screenshots* directory of the repository associated to this report [1].

## 4.1 Top Panel

The code for the Top Panel is located in */Panels/TopPanel*, in the files *Preset.cpp* and *Preset.h*. As it has been previously anticipated, the Top Panel manages presets (class *StatePresets*) and how the user can interact with them via the GUI (class *StateComponent*); this includes both user presets and temporary presets for doing A/B comparisons.

Presets are saved in a *.fire* format, which is actually just a XML file containing a single elements (*WINGSFIRE*) with all the effect parameters as its attributes. For this reason, the first part of *Preset.cpp* is dedicated to handling XML files. In this part the class *StateAB* is also implemented. Presets for A/B comparisons are treated as normal

presets, just saved in a Temp folder. A `PresetNameSorter` class, whose job is to sort the presets in the UI, is also implemented. Interestingly, the sorter takes a generic `const juce::String &attributeToSortBy` as its input and sorts based on it; this means that multiple sorting criteria could be implemented, though the author only used the preset names. A second parameter defines whether to invert the order of the elements or not, but this is also not used and always set to true (do not invert).

The second part of the file (lines 208-433) is dedicated to the `StatePresets` class. It handles loading, saving and displaying presets in the Menu.

The third part of the file (line 438-837) is dedicated to the `StateComponent` class which actually implements all functionalities in the GUI. This means an A/B button, a "Copy" button for AB comparisons, a preset menu, where presets are listed in alphabetic order, divided in folders, then a "Save" button and finally the "Menu" button, which can be used to load the default preset and open or rescan the preset folder and its subfolders.

## 4.2 Spectrogram Panel

The Spectrogram Panel part is dedicated to the part of the GUI when the Spectrogram panel is shown. It is highly integrated with the Control Panel, in fact it is possible to use this part of the GUI to create up to four bands, control their boundaries and eventually turn bypass the effect in them if needed.

We will not comment this part of the code any further since it goes out of the scope of this report.

## 4.3 Control Panel

Four maximum bands means four times the controls to implement, and in fact that's what happens in the Control Panel part. This part of the code is responsible for "spawning" controls in the GUI, both for individual bands and for the whole plugin (global controls); the code for the four visualisers (distortion input-output response, oscilloscope, meter and a polar vectorscope) in the bottom left corner is also included in a special folder in this section.

We'll focus on the oscilloscope as an example and we'll omit the rest in order not to overflow this report with extra informations.

The general idea is to get the data from a buffer calling `processor.getHistoryArrayL` (also repeating it for the right channel with `getHistoryArrayR` if the signal is stereo), find the maximum amplitude `maxValue`; if it's bigger than a minimum (in order to avoid divisions by very low number or even zeros), normalise the values in the array and then do a for loop to draw the waveform using the `strokePath` function of a `juce::Graphics` object received as an input. The function uses a scaling factor on the for loop index in order to extract amplitude values from a smaller chunk of samples distributed uniformly in the array. Since `strokePath` requires a `juce::Path` object as an input, one is created (two for stereo) and is (are) constantly updated with each iteration using values from the extracted samples.

## 5 Conclusion

Thanks to the correct implementation of the overdrive processor and the many waveshaping functions the *Fire* is a viable open-source option to the many multi-band distortion plugins on the market. It does its job with great results and the intuitive graphical interface makes it easy to use. A plus point of the code is that it is written following a modular approach, such that the whole audio process can be seen in an high level by a main function while the actual work is divided between low-level blocks. At the same time it could benefit from a code cleanup (as we already said there are files like *diodeWDF.h* and *WDF.h*, but also *Delay.h* that are practically not used) and from some commenting - some of the most important files are not commented, or comments are not explanations but rather TODOs or notes by the author. Another future improvement could be more flexibility in the *Global Effect* part, particularly giving the user the freedom to have more than 3 filters and decide for each the filtering type.

## References

- [1] *Analysis of the Fire plugin - GitHub repository*. URL: <https://github.com/polimi-cmls-23/group9-hw-JUCE-TheSineOfTheTimes>.